

# A4\_Report\_123090612

## Q1 Mirror

### ①What are the possible solutions for the problem?

#### 1. Standard Shortest Path (Dijkstra's Algorithm):

**When to use:** When there are **no mandatory edges** to traverse, or the mandatory edges are implicitly covered by the starting and ending nodes.

**Approach:**

- Use **Dijkstra's algorithm** to find the shortest path between the start node (  $s$  ) and the target node (  $t$  ).
- This is the simplest solution, where we don't worry about any constraints related to edge traversal.

#### 2. Modified Dijkstra's Algorithm with Bitmasking:

**When to use:** When **mandatory edges** must be traversed as part of the shortest path.

**Approach:**

- **Bitmasking** is used to represent which mandatory edges have been traversed.
- For each query, we run a **modified Dijkstra's algorithm**, where each state (node) is associated with a bitmask that tracks the mandatory edges.
- If the destination node (  $t$  ) is reached and all mandatory edges have been visited (i.e., the bitmask matches the required value), we print the shortest time.

#### 3. Shortest Path with Constraints (Multi-Path Search):

**When to use:** For larger graphs where multiple queries are involved, and the mandatory edges vary for each query.

**Approach:**

- Precompute the **shortest path** from each node to every other node using **Floyd-Warshall** or **Dijkstra** (for each node as the source).
- For each query, combine the precomputed shortest paths and the required mandatory edges.
- This solution avoids recalculating the shortest paths from scratch for every query.

### ②How do you solve this problem?

#### 1. Helper Functions for Priority Queue

This section defines a set of functions to implement a basic priority queue using a list-based heap.

```
def heap_push(heap, item):
    heap.append(item)
    sift_up(heap, len(heap) - 1)
def heap_pop(heap):
    if not heap:
        return None
    top = heap[0]
    last = heap.pop()
    if heap:
        heap[0] = last
        sift_down(heap, 0)
    return top
def sift_up(heap, idx):
    while idx > 0:
        parent = (idx - 1) // 2
        if heap[parent][0] > heap[idx][0]:
            heap[parent], heap[idx] = heap[idx], heap[parent]
            idx = parent
        else:
            break
def sift_down(heap, idx):
    n = len(heap)
    while True:
        left = 2 * idx + 1
        right = 2 * idx + 2
        smallest = idx
        if left < n and heap[left][0] < heap[smallest][0]:
            smallest = left
        if right < n and heap[right][0] < heap[smallest][0]:
            smallest = right
        if smallest != idx:
            heap[smallest], heap[idx] = heap[idx], heap[smallest]
            idx = smallest
        else:
            break
```

- **Heap Push:** `heap_push` adds a new element to the heap and ensures the heap property is maintained by "sifting" the new element upwards if necessary ( `sift_up` ).
- **Heap Pop:** `heap_pop` removes the top element from the heap, which is the smallest one due to the nature of the heap. The heap then re-adjusts by moving the last element to the top and "sifting it down" ( `sift_down` ) to maintain the heap property.

- **Sifting:** The `sift_up` and `sift_down` functions ensure the heap maintains the min-heap property (smallest element is always at the root).

## 2. Input Reading Functions

This section contains functions that handle input reading. These are specifically written to read multiple lines of integers from the input.

- `readints` : This function reads a line of input, strips it of extra spaces, splits it into integers, and returns the list. If the line is empty, it continues to read.
- `read_edge` : This is similar to `readints` , but specifically designed for reading edges, ensuring that the input is properly formatted as three integers per line.

## 3. Main Logic

### Initial Setup

```
def main():
    first_line = readints()
    while not first_line:
        first_line = readints()
    n, m, q = first_line
```

- **Reading Input:** The first line contains three integers: `n` (the number of nodes), `m` (the number of edges), and `q` (the number of queries). This is stored in `n` , `m` , and `q` respectively.

### Reading Edges and Building the Graph

```
edges = [None] # 1-based indexing
graph = [[] for _ in range(n + 1)]
for edge_idx in range(1, m + 1):
    edge = read_edge()
    while not edge:
        edge = read_edge()
    u, v, w = edge
    edges.append((u, v, w))
    graph[u].append((v, w, edge_idx))
    graph[v].append((u, w, edge_idx))
```

- `edges` : A list to store all the edges. It starts with `None` for 1-based indexing.
- `graph` : An adjacency list representation of the graph. Each node has a list of tuples representing neighboring nodes, the weight of the edge, and the edge index.
- **Edges:** The code loops through the edges and adds each one to the `edges` list and the adjacency list `graph` .

## 4. Reading Queries

```
queries_E = []
for _ in range(q):
    ki_list = readints()
    while not ki_list:
        ki_list = readints()
    k_i = ki_list[0]
    E_j = []
    while len(E_j) < k_i:
        e_list = readints()
        if not e_list:
            continue
        E_j += e_list
    queries_E.append(E_j[:k_i])
```

- `queries_E` : This list will store the sets of edges that must be traversed for each query.
- **Reading Edges for Each Query:** For each query, it reads the list of edge indices `E_j` that Lee must pass through. The number of edges to traverse (`k_i` ) is read first.

### Reading Start and End Nodes for Queries

```
queries_ST = []
for _ in range(q):
    st_list = readints()
    while not st_list:
        st_list = readints()
    s_j, t_j = st_list
    queries_ST.append((s_j, t_j))
```

- `queries_ST` : This list stores the start (`s_j` ) and end (`t_j` ) nodes for each query.

## 5. Processing Each Query

### Dijkstra's Algorithm for Shortest Path

```
for query_idx in range(q):
    E_j = queries_E[query_idx]
    s_j, t_j = queries_ST[query_idx]
```

```

k_i = len(E_j)
if k_i == 0:
    # If there are no edges to traverse, simply find the shortest path
    # using standard Dijkstra's algorithm without mask

```

- **Processing Queries:** For each query, it checks if there are specific edges that must be passed through. If there are no such edges, it directly finds the shortest path between  $s_j$  and  $t_j$  using a standard Dijkstra's algorithm.

### Dijkstra's Algorithm (With Masks for Specific Edges)

If there are edges to pass through, the code applies a modified version of Dijkstra's algorithm that tracks which mandatory edges have been visited using a bitmask.

```

# Create a mapping from edge_index to bit position
E_j_map = {}
for bit_pos in range(k_i):
    edge_index = E_j[bit_pos]
    E_j_map[edge_index] = bit_pos
target_mask = (1 << k_i) - 1

```

- **Bitmasking:** This step maps each edge in  $E_j$  to a unique bit position. The  $target\_mask$  represents the bitmask where all mandatory edges have been traversed.

```

# Initialize Dijkstra
heap = []
heap_push(heap, (0, s_j, 0))
visited = [set() for _ in range(n + 1)]
answer_found = False
while heap:
    current = heap_pop(heap)
    if current is None:
        break
    current_time, current_node, mask = current
    if current_node == t_j and mask == target_mask:
        print(current_time)
        answer_found = True
        break

```

- **Modified Dijkstra:** Here, the heap stores the current time, node, and the bitmask. If the destination node  $t_j$  is reached and all required edges have been passed (i.e., the bitmask matches  $target\_mask$ ), the shortest time is printed.

```

if mask in visited[current_node]:
    continue
visited[current_node].add(mask)
for neighbor in graph[current_node]:
    v, w_edge, edge_idx = neighbor
    new_mask = mask
    if edge_idx in E_j_map:
        bit = E_j_map[edge_idx]
        new_mask = mask | (1 << bit)
    heap_push(heap, (current_time + w_edge, v, new_mask))
if not answer_found:
    print(-1)

```

- **Visiting Nodes:** The algorithm checks whether a node with a particular mask has already been visited to avoid unnecessary recalculations. It then updates the mask for each edge and continues the process.

## 6. Final Output

### ③Why is your solution better than others?

My solution is **better** because it efficiently handles multiple queries with varied edge requirements, scales well with large graphs, and guarantees optimal results using a combination of **Dijkstra's algorithm** and **bitmasking**. This combination ensures that the solution is both **time-efficient** and **correct**, making it the best choice for the problem.

## Q2 Violet

### ①What are the possible solutions for the problem?

#### 1. Binary Search on Bottleneck Value with BFS/DFS

This solution is the one implemented in the provided code. It combines **binary search** with **BFS** (or DFS) to find the maximum possible minimum number of flowers along a path between two flowerbeds.

- Use **binary search** over the edge flower values (from 0 to the maximum edge weight) to find the maximum value of the minimum flowers along a path.
- For each mid-point in the binary search, use **BFS/DFS** to check if there is a path between  $s_i$  and  $t_i$  where every edge on the path has at least  $mid$  flowers.
- If such a path exists, increase the lower bound of the binary search; otherwise, decrease the upper bound.

#### 2. Union-Find (Disjoint Set Union - DSU) with Path Compression and Union by Rank

Another possible solution could use **Union-Find** (or DSU) to dynamically maintain connected components as the graph changes. This approach focuses on finding connected components efficiently and updating them as edges change.

- Use a **Union-Find** data structure to keep track of the connected components, where each component represents a group of flowerbeds that are connected.
- For each edge update, adjust the Union-Find structure by merging sets (flowerbed groups) as necessary.
- For each query, find the connected components for flowerbeds `s_i` and `t_i`. If they belong to the same component, we can compute the maximum bottleneck value in that component.

### 3. Dynamic Shortest Path Algorithms

If the goal is to consider all possible paths between `s_i` and `t_i` and dynamically change the graph, we might consider **dynamic shortest path algorithms** like **Dijkstra's algorithm** or **Bellman-Ford**. These algorithms can be adapted to compute the minimum bottleneck on a path.

- Use Dijkstra's or Bellman-Ford to find the shortest or the maximum bottleneck path from `s_i` to `t_i`.
- Each time an edge weight changes, re-run the algorithm to update the shortest or maximum bottleneck path.

## ②How do you solve this problem?

### 1. Input Parsing

- The first part reads the number of **flowerbeds** `n` and the number of **edges** `m`.
- Then, we read the edges. Each edge is defined by two nodes `u` and `v` (the flowerbeds it connects) and `w` (the number of flowers on that edge).
- The code also ensures that the flowerbed indices `u` and `v` are ordered in a consistent way (`u < v`). This is done to make sure the graph edges are stored in a uniform order, which helps later in searching for specific edges.
- The list `edge_list` is then sorted lexicographically by `u` and `v`, which will be useful for binary searching edges later.

### 2. Graph Representation

- An **adjacency list** `adj` is created to represent the graph. Each flowerbed (node) will store a list of tuples, where each tuple represents a neighboring flowerbed and the index of the edge connecting them.
- For every edge in `edge_list`, both flowerbeds `u` and `v` are updated in the adjacency list. This ensures that the graph is undirected, i.e., if there's an edge between flowerbed `u` and flowerbed `v`, both `u` will point to `v` and `v` will point to `u`.

### 3. Reading Change Requests

- This section handles the reading of the **changes** that occur over time.
- `q` is the number of changes (days), and for each change, we first read `k_i`, the number of paths affected by the change.
- Each change updates the number of flowers on certain paths (edges). The path between flowerbeds `a` and `b` will now have `c` flowers instead of its previous count.
- The list `changes` stores all the changes for each day in the format `[(a, b, c), ...]`.

### 4. Reading Queries

- After reading the changes, we then read the **queries**.
- For each query (which is asked for each day, including day `0` before any changes), we need to find the path from flowerbed `s_i` to `t_i` that maximizes the minimum number of flowers along the path.
- These queries are stored in the list `queries`.

### 5. Helper Functions

#### Edge Index Lookup

```
# Function to find the index of an edge in the sorted edge_list
def find_edge_index(a, b):
    if a > b:
        a, b = b, a
    l = 0
    r = len(edge_list) - 1
    while l <= r:
        mid = (l + r) // 2
        if edge_list[mid][0] == a and edge_list[mid][1] == b:
            return mid
        elif edge_list[mid][0] < a or (edge_list[mid][0] == a and edge_list[mid][1] < b):
            l = mid + 1
        else:
            r = mid - 1
    return -1 # Edge not found
```

- This function finds the index of an edge in the sorted `edge_list`.
- The binary search approach is used to efficiently locate the edge `(a, b)` by comparing flowerbed indices, ensuring the edge exists in the list and is correctly positioned.

#### Connectivity Check

```
# Function to check if there's a path from s to t with all edges having at least w flowers
def is_connected(s, t, w):
    if s == t:
        return True
    visited = [False] * (n + 1)
    queue = [s]
```

```

visited[s] = True
head = 0
while head < len(queue):
    u = queue[head]
    head += 1
    for (v, edge_idx) in adj[u]:
        if not visited[v] and edge_list[edge_idx][2] >= w:
            if v == t:
                return True
            visited[v] = True
            queue.append(v)
return False

```

- This function checks if there exists a **path** from `s` to `t` such that every edge along the path has at least `w` flowers.
- It uses a **BFS (breadth-first search)** approach to explore the graph. If all edges encountered in the path have at least `w` flowers, the function returns `True`, meaning there's a valid path from `s` to `t` with this constraint. If no such path exists, it returns `False`.

### Binary Search for Maximum Bottleneck

```

# Function to find the maximum bottleneck value between s and t using binary search
def find_max_bottleneck(s, t):
    low = 0
    high = max(edge[2] for edge in edge_list) if edge_list else 0
    ans = 0
    while low <= high:
        mid = (low + high) // 2
        if is_connected(s, t, mid):
            ans = mid
            low = mid + 1
        else:
            high = mid - 1
    return ans

```

- This function uses **binary search** to find the **maximum minimum number of flowers** that can be encountered along a path from `s` to `t`.
- The search is performed on the range of possible flower counts (`w`). For each value `w`, the `is_connected()` function is used to check if there exists a valid path from `s` to `t` with all edges having at least `w` flowers.
- The binary search tries to maximize `w`, and the result is the maximum possible minimum number of flowers for the path.

## 6. Main Logic

```

# Process each day and handle the changes
for day in range(q + 1):
    if day > 0:
        # Apply the (day-1)th change
        change = changes[day - 1]
        for (a, b, c) in change:
            idx = find_edge_index(a, b)
            if idx != -1:
                edge_list[idx][2] = c
    # Handle the query for the current day
    s, t = queries[day]
    res = find_max_bottleneck(s, t)
    print(res)

```

- This loop iterates over all days (`q + 1` days).
- For each day:
  - If it's not the first day, the program applies the changes from the previous day to the graph by updating the flower counts of affected edges.
  - It then solves the query for that day by calling `find_max_bottleneck(s, t)` to get the maximum minimum number of flowers for the path between `s` and `t`.
  - The result for each query is printed.

## ③Why is your solution better than others?

### 1. Efficient Handling of Dynamic Changes

- The problem involves **dynamic updates** to the graph (edge weights change over time). The binary search approach is well-suited for this because it efficiently adjusts the graph's state each day by applying the changes and then recalculating the query result.
- This is much more efficient than algorithms like **Max-Flow** or **Dijkstra's** which might require full recomputation of paths after every update, leading to higher time complexity.

### 2. Optimal Search with Binary Search

- By using **binary search** on the possible edge weights, the solution narrows down the maximum possible "bottleneck" (minimum flowers on any edge in the path) in a logarithmic manner. This dramatically reduces the number of checks (from potentially testing all edge weights to just  $\log(W)$  where  $W$  is the maximum flower count).
- This is far more efficient than methods like **Dijkstra's algorithm** or **Max-Flow** that would need to explore the entire graph for each query.